

At the beginning of the task where we are asked to initialize parameters with He initialization, I was a little bit confused on what is the purpose of using random initialization and zero initialization for weight matrices and biases respectively. After doing some research, I found out that randomness is used to find a good enough set of weights for the specific mapping function. Moving forward, I encountered another issue when trying to implement the Sigmoid function. In the original function, there exists a condition such that when $Z \geq 0$, then $\text{Sigmoid}(Z) = 1/(1+\exp(z))$, otherwise, $\text{Sigmoid}(Z) = \exp(Z)/(1+\exp(Z))$. When implementing this to my code, I did

if $Z \geq 0$:

$A = 1/(1 + \exp(-Z))$

cache = Z

else:

$A = \exp(Z)/(1 + \exp(Z))$

cache = Z

```
tmp = []
for items in Z[0]:
    if items >= 0:
        tmp.append(1/(1 + np.exp(-items)))
        cache = Z
    else:
        tmp.append(np.exp(items)/(1 + np.exp(items)))
        cache = Z
A = np.array([tmp])
```

However, this yields an error since I just Z is not a number, but rather an array. Hence, I decided to iterate it.

Next, I also encountered a small problem when computing the Binary cross-entropy loss. I used np.sum and np.multiply to calculate the cost at first, and it computed the correct result but returns a wrong data type (numpy.float64 instead of numpy.ndarray). I figured that this must have something to do with the way I calculate the variable cost, and I noticed that I should use np.dot instead of np.multiply and np.sum to get numpy.ndarray as the data type.

BINARY CLASSIFIER IMPLEMENTATION

We are implementing the function `L_layer_model` to build the Binary Classifier. In this function, we first need to initialize the parameters by calling the function `initialize_parameters_deep` and pass `layers_dims`, which is the list containing the input size and each layer size as the parameter. We will get a dictionary with our parameters W_{l-1} and b_{l-1} .

Next, inside the for loop which iterates for a number of `num_iterations` times, we want to get the last post-activation value and a list of caches, and we can do so by calling the `L_model_forward(X, parameters, classes)` function. Afterwards, we can compute the Binary cross-entropy loss by calling the implemented `compute_BCE_cost(AL, Y)`. Then, we can call `L_model_backward(AL, Y, caches, classes)` for the backward propagation which will return dictionary with gradients, which will be used when we update the parameters. Finally, we can call `update_parameters(parameters, grads, learning_rate)` to update the parameters using gradient descent on every W_l and b_l for $l = 1, 2, \dots, L$.

For the hyperparameter tuning, I used [4, 1] as the value of `layer_dims`, 0.0075 as the learning rate, and 3000 as the number of iterations. In the end, my accuracy value is 1.0.

```
layers_dims = [4, 1]
parameters = L_layer_model(X_train, y_train, layers_dims, learning_rate = 0.0075, num_iterations = 3000, print_cost = True, classes=2)
```

FUNCTION IMPLEMENTATIONS

To be able to implement the Binary Classifier, there are a few supporting functions that we need to implement beforehand.

initialize_parameters(n_x, n_y):

Calculate the parameters W_1 and b_1 using n_x and n_y , which are the size of input and output layer respectively. We use random initialization and multiply it by $\sqrt{2./n_x}$ for W_1 and zero initialization and multiply it with $\sqrt{2./n_y}$ for b_1 .

initialize_parameters_deep(layer_dims):

Given the array with the dimensions of each layer as the values, we can use this to compute the parameters W_{l-1} and b_{l-1} using random and zero initialization and multiply it by $\sqrt{2./\text{layer_dims}[l-1]}$ where $\text{layer_dims}[l-1]$ represents the dimension of the previous layer, since dimension of current layer is represented by $\text{layer_dims}[l]$.

linear_forward(A, W, b):

Here, we just need to calculate Z, the pre-activation parameter, which is the dot product of W and A and then summed with b.

sigmoid(Z):

Calculates the activation value using the given formula:
$$\text{Sigmoid: } \sigma(Z) = \begin{cases} \frac{1}{1+e^{-Z}}, & \text{if } Z \geq 0 \\ \frac{e^Z}{1+e^Z}, & \text{otherwise} \end{cases}$$

relu(Z):

Relu(Z) returns the greater value between Z and 0.

linear_activation_forward(A_prev, W, b, activation):

In this function, we just need to call `linear_forward(A_prev, W, b)` and store it to Z and `linear_cache`, and depending on the value stored in the string `activation`, we call the respective function, either `sigmoid` or `relu`, and pass Z as the parameter.

compute_BCE_cost(AL, Y):

Calculate the cost using the given formula:

$$-\frac{1}{m} \sum_{i=1}^m (y^{(i)} \log(a^{[L](i)}) + (1 - y^{(i)}) \log(1 - a^{[L](i)}))$$
$$dW^{[l]} = \frac{\partial \mathcal{J}}{\partial W^{[l]}} = \frac{1}{m} dZ^{[l]} A^{[l-1]T}$$
$$db^{[l]} = \frac{\partial \mathcal{J}}{\partial b^{[l]}} = \frac{1}{m} \sum_{i=1}^m dZ^{[l](i)}$$
$$dA^{[l-1]} = \frac{\partial \mathcal{L}}{\partial A^{[l-1]}} = W^{[l]T} dZ^{[l]}$$

linear_backward(dZ, cache):

Calculate dW, db, and dA_prev using the given formulas:

sigmoid_backward(dA, cache):

Calculate dZ using the formula $dZ^{[l]} = dA^{[l]} * g'(Z^{[l]})$, where $g'(Z^{[l]})$ can be found by counting the derivative of the sigmoid function using the formula: $\sigma'(Z^{[l]}) = \sigma(Z^{[l]})(1 - \sigma(Z^{[l]}))$. This can be implied in Python code by:

```
dZ = dA * ((1/(1+np.exp(-Z))) * (1-(1/(1+np.exp(-Z)))))
```

relu_backward(dA, cache):

First convert dA to `numpy.array` and set dZ to 0 when $Z \leq 0$.

linear_activation_backward(dA, cache, activation):

Depends on the value of the string `activation` (either “relu” or “sigmoid”), call `relu_backward` or `sigmoid_backward` to find the value of dZ which we will use to find the value of dA_prev, dW, and db by calling `linear_backward(dZ, linear_cache)`.

update_parameters(parameters, grads, learning_rate):

Here we just need to update the parameters W_l and b_l for $l = 1, 2, \dots, L$ using gradient descent:
$$W^{[l]} = W^{[l]} - \alpha dW^{[l]}$$
$$b^{[l]} = b^{[l]} - \alpha db^{[l]}$$
 where α is the learning rate. This can be done in Python using:

```
for l in range(L):
    parameters["W" + str(l+1)] = parameters["W" + str(l+1)] - learning_rate * grads["dw" + str(l+1)]
    parameters["b" + str(l+1)] = parameters["b" + str(l+1)] - learning_rate * grads["db" + str(l+1)]
```